


Dhwani Kavach - full-scale evaluation

ASVspoof 2019 LA eval split, AUDDT clip protocol (4 s window), scored by the live model path


reported	value
ROC-AUC	0.9489
EER	88.40 %
TPR at 1% FPR	76.94 %
clips scored	71,237
bonafide / spoof	7,355 / 63,882
model file	model.pt a7aeb383a12b
parameters	236,141
accuracy @ tau=0.50	82.64 %
accuracy @ theta_lo=0.62	74.43 %
recall / FPR at theta_lo	71.53 % / 0.38 %
logit bias (monotone)	-3.677
calibration fitted on	29 voiced windows, 2026-09-09
robustness rows	10 conditions
policy: real warned / spoof caught (dwell off)	0.13 % / 58.53 %
policy: spoofs caught	0.0 %
git / torch	d4c66e3 / 2.14.0+cpu

This table was written by 18_evaluate_auddt.py in the same run that produced metrics.json.
No number on this page is typed by hand.


Detection error trade-off, full eval split




Hardest attack families first



Robustness: what noise, codecs and short clips do to us (1200 clips per condition)



Ranking quality and the live operating point, per condition



How this was measured

split / protocol	: eval, one 4 s window per clip, seed 20260925
features	: log-Mel 80 / 16 kHz / fft 512 / hop 160 / power 2.0 / ref 1.0 then per-window standardise (mean, std) - identical to demo_server
model	: MelCNN, results\cnn_v2\model.pt, 236,141 params
decode	: soundfile + librosa resample (no torchaudio/FFmpeg in the scoring path)
codec conditions	: real ffmpeg (imageio-ffmpeg) mulaw, alaw, amrnb, opus16, mp3_64
padding	: short clips are tiled to 4 s by fix_length (product behaviour)

Read the limits:

- * 90% of the scored clips are bonafide, so accuracy alone is misleading.
- * EER here is a clip-level number; the product's decision is the streamed policy table.
- * the logit bias only moves the operating point (monotone): AUC/EER are unchanged by it, which is why the clean and calibrated rows above agree.
- * gates printed by the run: model-file=PASS, calibration=PASS, parity-vs-live-server=PASS, manifest=PASS,

artifacts: metrics.json, auddt_evaluation_report.md, auddt_evaluation_report.pdf
git d4c66e3 | torch 2.14.0+cpu | Windows-10-10.0.26200-SP0 | 2026-09-25 10:16:38